

assignment11

November 12, 2023

```
[1]: import pandas as pd
import numpy as np
from transformers import BertTokenizer, BertForSequenceClassification
import torch
from sklearn.model_selection import train_test_split
from torch.utils.data import DataLoader, RandomSampler, SequentialSampler, TensorDataset

[2]: file_path = '/Users/pan/Desktop/dataanalytics/bodybuilding_nutrition_products.csv'
df = pd.read_csv(file_path)

[3]: threshold_for_favorite = 100

# Create the binary target column
df['is_favorite'] = np.where(df['number_of_reviews'] >= threshold_for_favorite, 1, 0)

[4]: tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

input_ids = []
attention_masks = []

# Assuming 'Recipe_name' is the text column
for sentence in df['brand_name']:
    encoded_dict = tokenizer.encode_plus(
        sentence,
        add_special_tokens = True,
        max_length = 64,
        pad_to_max_length = True,
        return_attention_mask = True,
        return_tensors = 'pt',
    )

    input_ids.append(encoded_dict['input_ids'])
    attention_masks.append(encoded_dict['attention_mask'])
```

```
input_ids = torch.cat(input_ids, dim=0)
attention_masks = torch.cat(attention_masks, dim=0)
labels = torch.tensor(df['is_favorite'])
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

/Users/pan/anaconda3/lib/python3.11/site-

packages/transformers/tokenization_utils_base.py:2418: FutureWarning: The `pad_to_max_length` argument is deprecated and will be removed in a future version, use `padding=True` or `padding='longest'` to pad to the longest sequence in the batch, or use `padding='max_length'` to pad to a max length. In this case, you can give a specific length with `max_length` (e.g. `max_length=45`) or leave max_length to None to pad to the maximal input size of the model (e.g. 512 for Bert).

```
warnings.warn(
```

```
[5]: train_inputs, validation_inputs, train_labels, validation_labels =
      ↪train_test_split(input_ids, labels, random_state=2018, test_size=0.1)
train_masks, validation_masks, _, _ = train_test_split(attention_masks, labels,
      ↪random_state=2018, test_size=0.1)

train_data = TensorDataset(train_inputs, train_masks, train_labels)
train_sampler = RandomSampler(train_data)
train_dataloader = DataLoader(train_data, sampler=train_sampler, batch_size=32)

validation_data = TensorDataset(validation_inputs, validation_masks,
      ↪validation_labels)
validation_sampler = SequentialSampler(validation_data)
validation_dataloader = DataLoader(validation_data, sampler=validation_sampler,
      ↪batch_size=32)
```

```
[6]: model = BertForSequenceClassification.from_pretrained(
      "bert-base-uncased",
      num_labels = 2,
      output_attentions = False,
      output_hidden_states = False,
    )
```

Some weights of BertForSequenceClassification were not initialized from the model checkpoint at bert-base-uncased and are newly initialized:

```
['classifier.bias', 'classifier.weight']
```

You should probably TRAIN this model on a down-stream task to be able to use it for predictions and inference.

```
[7]: import pandas as pd
import numpy as np
import torch
from transformers import BertTokenizer, BertForSequenceClassification, AdamW, \
    get_linear_schedule_with_warmup
from sklearn.model_selection import train_test_split
from torch.utils.data import DataLoader, RandomSampler, SequentialSampler, \
    TensorDataset
```

2023-11-12 15:15:33.524143: I tensorflow/core/platform/cpu_feature_guard.cc:182] This TensorFlow binary is optimized to use available CPU instructions in performance-critical operations.
To enable the following instructions: AVX2 FMA, in other operations, rebuild TensorFlow with the appropriate compiler flags.

```
[9]: threshold_for_favorite = 100

# Create the binary target column
df['is_favorite'] = np.where(df['number_of_reviews'] >= threshold_for_favorite, \
    1, 0)
```

```
[10]: tokenizer = BertTokenizer.from_pretrained('bert-base-uncased')

input_ids = []
attention_masks = []

# Assuming 'Recipe_name' is the text column
for sentence in df['brand_name']:
    encoded_dict = tokenizer.encode_plus(
        sentence,
        add_special_tokens = True,
        max_length = 64,
        pad_to_max_length = True,
        return_attention_mask = True,
        return_tensors = 'pt',
    )

    input_ids.append(encoded_dict['input_ids'])
    attention_masks.append(encoded_dict['attention_mask'])

input_ids = torch.cat(input_ids, dim=0)
attention_masks = torch.cat(attention_masks, dim=0)
labels = torch.tensor(df['is_favorite'])
```

Truncation was not explicitly activated but `max_length` is provided a specific value, please use `truncation=True` to explicitly truncate examples to max length. Defaulting to 'longest_first' truncation strategy. If you encode pairs

of sequences (GLUE-style) with the tokenizer you can select this strategy more precisely by providing a specific strategy to `truncation`.

```
[11]: # Creating DataLoaders for Training and Validation:
# Split data into train and validation sets
train_inputs, validation_inputs, train_labels, validation_labels = \
    ↪train_test_split(input_ids, labels, random_state=2018, test_size=0.1)
train_masks, validation_masks, _, _ = train_test_split(attention_masks, labels, \
    ↪random_state=2018, test_size=0.1)

# Create the DataLoader for our training set
train_data = TensorDataset(train_inputs, train_masks, train_labels)
train_sampler = RandomSampler(train_data)
train_dataloader = DataLoader(train_data, sampler=train_sampler, batch_size=32)

# Create the DataLoader for our validation set
validation_data = TensorDataset(validation_inputs, validation_masks, \
    ↪validation_labels)
validation_sampler = SequentialSampler(validation_data)
validation_dataloader = DataLoader(validation_data, sampler=validation_sampler, \
    ↪batch_size=32)
```

```
[12]: # Setting Up Optimizer and Scheduler:
# Note: AdamW is a class from the huggingface library (as opposed to pytorch)
optimizer = AdamW(model.parameters(),
                  lr = 2e-5, # args.learning_rate - default is 5e-5
                  eps = 1e-8 # args.adam_epsilon - default is 1e-8
                  )

# Number of training epochs (authors recommend between 2 and 4)
epochs = 4

# Total number of training steps is [number of batches] x [number of epochs]
total_steps = len(train_dataloader) * epochs

# Create the learning rate scheduler
scheduler = get_linear_schedule_with_warmup(optimizer,
                                             num_warmup_steps = 0, # Default \
    ↪value in run_glue.py
                                             num_training_steps = total_steps)
```

```
/Users/pan/anaconda3/lib/python3.11/site-
packages/transformers/optimization.py:411: FutureWarning: This implementation of
AdamW is deprecated and will be removed in a future version. Use the PyTorch
implementation torch.optim.AdamW instead, or set `no_deprecation_warning=True`
to disable this warning
warnings.warn(
```

```
[13]: # Check if GPU is available and set the device accordingly
device = torch.device("cuda" if torch.cuda.is_available() else "cpu")
print(f"Using device: {device}")

# Also, make sure to move the model to the device
model.to(device)
```

Using device: cpu

```
[13]: BertForSequenceClassification(
  (bert): BertModel(
    (embeddings): BertEmbeddings(
      (word_embeddings): Embedding(30522, 768, padding_idx=0)
      (position_embeddings): Embedding(512, 768)
      (token_type_embeddings): Embedding(2, 768)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (encoder): BertEncoder(
      (layer): ModuleList(
        (0-11): 12 x BertLayer(
          (attention): BertAttention(
            (self): BertSelfAttention(
              (query): Linear(in_features=768, out_features=768, bias=True)
              (key): Linear(in_features=768, out_features=768, bias=True)
              (value): Linear(in_features=768, out_features=768, bias=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
            (output): BertSelfOutput(
              (dense): Linear(in_features=768, out_features=768, bias=True)
              (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
              (dropout): Dropout(p=0.1, inplace=False)
            )
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
)
(pooler): BertPooler(
```

```

        (dense): Linear(in_features=768, out_features=768, bias=True)
        (activation): Tanh()
    )
)
(dropout): Dropout(p=0.1, inplace=False)
(classifier): Linear(in_features=768, out_features=2, bias=True)
)

```

```

[14]: import time

def format_time(elapsed):
    '''
    Takes a time in seconds and returns a string hh:mm:ss
    '''
    # Round to the nearest second
    elapsed_rounded = int(round(elapsed))

    # Format as hh:mm:ss
    return str(datetime.timedelta(seconds=elapsed_rounded))

```

```

[ ]: import datetime
import random

# Training Loop:
# Seed value for reproducibility
seed_val = 42

random.seed(seed_val)
np.random.seed(seed_val)
torch.manual_seed(seed_val)
torch.cuda.manual_seed_all(seed_val)

# Store the average loss after each epoch so we can plot them
loss_values = []

for epoch_i in range(0, epochs):
    start_time = time.time()
    # Perform one full pass over the training set
    print('=====Epoch {:} / {:} ====='.format(epoch_i + 1, epochs))
    print('Training...')

    total_loss = 0
    model.train()

    # For each batch of training data...
    for step, batch in enumerate(train_dataloader):
        b_input_ids = batch[0].to(device)

```

```

b_input_mask = batch[1].to(device)
b_labels = batch[2].to(device)

# Clear any previously calculated gradients
model.zero_grad()

# Perform a forward pass (evaluate the model on this training batch)
outputs = model(b_input_ids, token_type_ids=None,
↪attention_mask=b_input_mask, labels=b_labels)
loss = outputs.loss
total_loss += loss.item()

# Perform a backward pass to calculate the gradients
loss.backward()

# Clip the norm of the gradients to 1.0 to prevent "exploding gradients"
torch.nn.utils.clip_grad_norm_(model.parameters(), 1.0)

# Update parameters and take a step using the computed gradient
optimizer.step()

# Update the learning rate
scheduler.step()

# Calculate the average loss over the training data
avg_train_loss = total_loss / len(train_dataloader)
loss_values.append(avg_train_loss)

time_elapsed = time.time() - start_time
print(" Average training loss: {0:.2f}".format(avg_train_loss))
print(" Training epoch took: {:}".format(format_time(time_elapsed)))

# Validation step
print("Running Validation...")

model.eval()

# Tracking variables
eval_loss, eval_accuracy = 0, 0
nb_eval_steps, nb_eval_examples = 0, 0
#

```

=====
Epoch 1 / 4 =====
Training...

```
[ ]: model = BertForSequenceClassification.from_pretrained("bert-base-uncased",
↪num_labels=2)
```

```
[ ]: output_dir = '/Users/yaoyaoliu/Documents/Graduate Class/2023/Fall 2023/293C/  
      ↪model_save/'  
      model.save_pretrained(output_dir)  
      tokenizer.save_pretrained(output_dir)
```

```
[ ]: model = BertForSequenceClassification.from_pretrained(output_dir)  
      tokenizer = BertTokenizer.from_pretrained(output_dir)
```

```
[ ]: # Classify New Sentences  
def classify_sentence(sentence):  
    # Tokenize the sentence  
    inputs = tokenizer.encode_plus(  
        sentence,  
        add_special_tokens=True,  
        max_length=64,  
        pad_to_max_length=True,  
        return_attention_mask=True,  
        return_tensors='pt',  
    )  
  
    # Move tensors to the same device as the model  
    input_ids = inputs['input_ids']  
    attention_mask = inputs['attention_mask']  
  
    # Get model predictions  
    with torch.no_grad():  
        outputs = model(input_ids, attention_mask=attention_mask)  
  
    # Convert output logits to softmax probabilities  
    probs = torch.nn.functional.softmax(outputs.logits, dim=1)  
  
    # Get the predicted class (the one with the highest probability)  
    predicted_class = torch.argmax(probs, dim=1).item()  
  
    return predicted_class  
  
# Example usage  
test_sentence = "The Best Vegan Breakfast Sandwich"  
prediction = classify_sentence(test_sentence)  
print(f"Predicted class for '{test_sentence}': {prediction}")
```

```
[ ]:
```